



Value Alignment in Intelligent Agents

Bad advice can be dangerous.



Bad action is worse.



Christopher Simpkins
Kennesaw State University
<https://chris.simpkins.phd/>

General Interests

- ▶ Multiagent systems
- ▶ Reinforcement learning
- ▶ Deep learning
- ▶ Causality
- ▶ ...

Current thrust: MARL approaches to value alignment