



Value Alignment as a Theory of Mind Problem



Goal: learn values from behavior.

- ▶ Not just “don’t run over pedestrians” but “don’t harm humans.”
- ▶ Start with action restrictions
- ▶ Learning approach
 - ▶ Don’t replace the Genie Problem with the knowledge acquisition bottleneck/bitter lesson.

Advances needed:

- ▶ Machine theory of mind
- ▶ Imitation/apprenticeship learning
- ▶ Offline reinforcement learning
- ▶ World models – causality

Impacts:

- ▶ Autonomous multi-agent systems
- ▶ Human-AI teaming
- ▶ Safe autonomy